


Generation flow (merged variant) — Capabilities in one call for all Value Streams

A variant of the [generation flow](#) where **L2/L3 Capabilities** are produced by a **single batched call covering every Value Stream's stages**, instead of one call per Value Stream. Everything else is unchanged. This trades a little quality for fewer calls, lower input tokens, and simpler orchestration.

Generation flow (merged) — Capabilities in ONE call for all Value Streams



Why parallel — and the one place it's sequential

- Capabilities is now ONE batched call covering every Value Stream's stages — not one call per VS. Strict stage isolation + salvage keep each L3 under its true stage (0 mislink in the output).
- Business Needs stays one call per Value Stream; Description is independent of the stages.
- Sequential link: Business Needs and Capabilities both need the SELECTED STAGES, so they wait for Stage Selection. Call count drops from $3 + 2N$ to $4 + N$ (Capabilities $N \rightarrow 1$).

What changed

	per-VS (current)	merged (this variant)
Capabilities (L3)	one call per Value Stream	one call for ALL Value Streams
isolation	per-VS, stage isolation	strict stage isolation across all VS + salvage
everything else	—	identical

The merged call puts every Value Stream's selected stages — each with its own governed candidate L3 — into one prompt. Candidates stay per-stage (strict isolation), and a deterministic salvage step re-routes any L3 placed under the wrong stage to its true owner (a capability id belongs to exactly one stage), so the **output is correct regardless**.

Call count

For **N** approved Value Streams:

	per-VS	merged
Capabilities calls	N	1
total LLM calls	5 + 2N	6 + N
at N = 3	11	9
at N = 10	25	16

Capabilities collapses from N calls to 1, removing **N – 1** calls per ticket (9 at N = 10).

Quality — measured (merged L3 eval, 118 tickets)

metric	per-VS (one_call)	merged (all VS)
precision	0.53	0.43
recall	0.90	0.84
F1	0.67	0.57
cross-stage/VS mislink (raw)	5.1%	17.0%
mislink in the output (after salvage)	0	0
hallucinated ids	0.0%	0.0%

- **Recall drops ~6 points** (0.90 → 0.84) — the real cost; one big prompt spreads attention thinner.
- **Precision drops ~10 points** strict, but a pick-relevance judge showed only **~18% of the non-GT picks are genuinely irrelevant** — the rest is plausible capabilities the ground truth didn't tag (under-tagging), so the *real* precision gap is much smaller.
- **Mislink looks worse (17% vs 5% raw) but is 0 in the delivered output either way** — salvage fixes it. Nothing is hallucinated.

Cost (GPT-5-mini: \$0.25/1M in, \$2.00/1M out; average tokens)

Capabilities only:

	per-VS	merged
calls	10	1
input tokens	58,470	~14,500
output tokens	6,990	6,990
capabilities cost / ticket (10 VS)	\$0.0286	\$0.0176

The **output tokens are the same** (you still generate the same L3, just in one response — that's the expensive \$2/1M side), so the saving is on **input**: the raw idea-card text is sent **once** instead of 10 times. **Capabilities cost drops ~38% (~\$0.011/ticket)** at average idea-card size — and **more on large idea cards**, where the per-call 24k raw text dominates (9 fewer copies of up to 24k tokens).

Per-ticket total (10 VS, average tokens): **~\$0.093 → ~\$0.082**.

Latency (measured)

	per-VS (one_call)	merged
unit	per Value Stream	per ticket (1 call)
avg / median	4.4s / 4.0s	10.9s / 6.4s
p90 / p95	—	21.5s / 34.6s
over 30s	—	5% of tickets (concurrency tail)

- Merged is **one ~6.4s call** vs N per-VS calls. If you run the per-VS calls **in parallel**, wall-clock is ~4.4s — so merged is slightly slower on the clock but **1 call instead of N**. If they ran sequentially, merged is far faster (6.4s vs 10×4.4s).
- The slow tail (≈5% of tickets over 30s, max 160s) is **concurrency-induced queueing**, not the model — big idea cards + concurrent load. Size timeouts for the p95 (~35s), not the average.

Verdict

Merged Capabilities is a **fair trade, not free**:

- Wins:** N → 1 calls (9 fewer at N=10), ~38% lower Capabilities input cost (more on big cards), simpler orchestration, less rate-limit exposure. 0 mislink in output, 0 hallucination.
- Cost:** ~6 pts recall, ~10 pts strict precision (much of it under-tagging, not real error), and a ~5% slow-call tail to size timeouts around.

Use merged when call-count / cost / rate-limits matter and the architect reviews the output anyway. **Keep per-VS** when maximum L3 recall is the priority. The same approach extends to **Business Needs** (a batched all-VS call exists — `--mode batched`); its quality numbers are pending.