

How much of the answer does our “similar past tickets” search find?

A plain-language report. We tested the search on 373 real tickets. For each one, the system pulls the most similar past tickets and shows them to the model as examples. This measures coverage: of a ticket's correct Value Streams, how many does the search actually surface.

"@6" means "when we show the top 6 pulled past tickets."

Bottom line


- The search **reliably finds the right Value Streams** — the correct ones show up for **90%** of a typical ticket's streams, and **96%** of tickets find at least one.
 - **78% of tickets find EVERY** correct Value Stream; **18%** find some but miss a few; only **4%** find none.
 - On average, a ticket has **3.2** correct Value Streams and the search finds **~2.9 of them** (90%).
 - The misses are concentrated in the **hard tickets** (those with several Value Streams) — there we find *most* streams (88%) but the *complete* set only **63%** of the time.
 - **Showing 6 past tickets is enough** — going to 10 adds only ~4 points of coverage.
 - **Bottom line:** the search puts the right answers in front of the model — **coverage is not the bottleneck.**
-

1. What happened to each ticket

When we show the 6 most similar past tickets, every ticket lands in exactly one of three groups (so they add up to 100%):

What happened to each ticket?

Every ticket is in exactly one group — the three add up to 100%.



How to read it. Each ticket has one or more correct Value Streams to find. Did the 6 examples contain them?

- **Found ALL — 78%:** every correct Value Stream was in the examples. Nothing missed. ✅
- **Found SOME — 18%:** found at least one, but missed a few. Usually the multi-stream tickets (chart 3).
- **Found NONE — 4%:** none of the correct streams showed up — the only real misses.

So **96% of tickets find at least something** (green + amber) and **78% find everything**.


How complete is each group? Even the "found some" group does well — they find about **two-thirds** of their streams on average:

Group	% of tickets	Of their own streams, how many found (avg)
Found ALL	78%	100%
Found SOME	18%	~68% (about 2 of every 3)
Found NONE	4%	0%

So the 18% "partial" tickets aren't finding just a sliver — they get most of the way and miss the last third (the long-tail streams on hard tickets).

One more number, measured differently: averaged across tickets, a typical ticket has **90%** of its correct streams present in the examples — in plain counts, **~2.9 of its 3.2** correct Value Streams are found on average. (That 90% is a per-ticket *average*, not a count of tickets — a ticket that finds 3 of its 4 streams contributes 75% to it. Both numbers say the same thing: coverage is strong.)

2. Not every ticket is equally hard




How to read it. Each bar is *how many of the 373 tickets* have that many correct Value Streams:

- **Just 1 — 193 tickets (52%):** one right answer to find. Easy.
- **2-4 — 98, 5-9 — 54, 10+ — 28:** the more streams a ticket has, the harder it is to find them *all*. The 28 tickets with 10+ streams (a few have 19) are the truly hard cases.

So "90% coverage" is blended across these — the 52% single-stream tickets are easy and lift the average.

3. Hard tickets: we find most, but rarely all



How to read it. Two ticket groups (**easy** = 1 stream, 193 tickets; **hard** = 2+, 180 tickets), each measured two ways:


- **Blue — "Found MOST streams (avg)":** the average fraction of a ticket's streams that were found. Easy 92%, hard 88% — nearly the same, so we find *most* of the answer even on hard tickets.
- **Green — "Found EVERY stream":** how often *all* streams were found. Easy 92%, hard only 63%.

In plain counts:

- **Easy ticket:** has 1 correct Value Stream → found 92% of the time (≈ **178 of 193** easy tickets get their one stream). With a single stream, "found most" and "found every" are the same — that's why the two bars match.
- **Hard ticket:** has ~5.6 correct Value Streams on average → the search finds ~4.9 of them (88%), but lands the *complete* set only 63% of the time (≈ **113 of 180** hard tickets).

So the hard-ticket gap (88% vs 63%) is the story: on a ~5.6-stream ticket we usually catch ~4.9 of them but **miss one or two from the long tail**, so we rarely get *all* of them.

4. Showing more tickets barely helps




How to read it. The x-axis is *how many past tickets we show* (6, 8, 10):

- **Green — Avg coverage:** climbs only 90% → 94% across the three ($\approx 2.9 \rightarrow 3.0$ of 3.2 streams found per ticket).
- **Blue — Found everything:** the share of tickets that get *all* their streams climbs only 78% → 85% ($\approx 291 \rightarrow 317$ of 373 tickets).

Each extra ticket adds a little coverage, but with **diminishing returns** — going from 6 to 10 examples finds only ~ 0.1 more stream on a typical ticket. **6 examples is enough** — more adds length without much benefit.

5. The right answer is usually at the very top



How to read it. For each ticket, the position of the *first* pulled ticket that carries a correct Value Stream:

- **Position 1 — 288 of 373 tickets (77%):** the very first result already carries a correct stream.
- **Positions 2, 3, 4–10 — ~76 tickets total:** a few need to look a little deeper.
- **None — 9 tickets (2%):** found nothing correct at all (the only true misses).

So the search not only *finds* the right streams, it usually *ranks them first* — for 3 in 4 tickets the top example is already on target.

All the numbers

Measure (per ticket)	Show 6	Show 8	Show 10
Streams found per ticket (of ~3.2 correct)	2.90	2.97	3.03
Avg coverage (fraction of streams found)	90%	92%	94%
Found at least one correct stream	96%	97%	98%
Found EVERY correct stream	78%	81%	85%

(A ticket has 3.2 correct Value Streams on average; the historic search surfaces ~2.9 of them. Note how the count barely moves — 2.90 → 3.03 — across 4 extra tickets: that's the diminishing returns from chart 4.)

Plain-language glossary

Term	What it means
Value Stream	The business category a ticket belongs to — what we predict.
Coverage	Of a ticket's correct Value Streams, how many showed up in the pulled examples.
Found everything	All of a ticket's correct Value Streams were found — none missed.
Easy / hard ticket	Easy = 1 correct Value Stream; hard = 2 or more.
@6 / @8 / @10	When showing the top 6 / 8 / 10 pulled past tickets.