

How much of the answer does our “similar past tickets” search find?

A plain-language report. We tested the search on **373 real tickets**. For each one, the system pulls the most similar past tickets and shows them to the model as examples. This measures **coverage**: of a ticket's correct **Value Streams**, how many does the search actually surface.

"@6" means "when we show the top 6 pulled past tickets."

Bottom line


- The search **reliably finds the right Value Streams** — the correct ones show up for **90%** of a typical ticket's streams, and **96%** of tickets find at least one.
 - **78% of tickets find EVERY** correct Value Stream; **18%** find some but miss a few; only **4%** find none.
 - On average, a ticket has **3.2** correct Value Streams and the search finds **~2.9 of them** (90%).
 - The misses are concentrated in the **hard tickets** (those with several Value Streams) — there we find *most* streams (88%) but the *complete* set only **63%** of the time.
 - **Showing 6 past tickets is enough** — going to 10 adds only ~4 points of coverage.
 - **Bottom line**: the search puts the right answers in front of the model — **coverage is not the bottleneck**.
-

1. What happened to each ticket

When we show the 6 most similar past tickets, every ticket lands in exactly one of three groups (so they add up to 100%):

What happened to each ticket?

Every ticket is in exactly one group — the three add up to 100%.



How to read it. Each ticket has one or more correct Value Streams to find. Did the 6 examples contain them?

- **Found ALL — 78%:** every correct Value Stream was in the examples. Nothing missed. ✅
- **Found SOME — 18%:** found at least one, but missed a few. Usually the multi-stream tickets (chart 3).
- **Found NONE — 4%:** none of the correct streams showed up — the only real misses.

So **96% of tickets find at least something** (green + amber) and **78% find everything**.


How complete is each group? Even the "found some" group does well — they find about **two-thirds** of their streams on average:

Group	% of tickets	Of their own streams, how many found (avg)
Found ALL	78%	100%
Found SOME	18%	~68% (about 2 of every 3)
Found NONE	4%	0%

So the 18% "partial" tickets aren't finding just a sliver — they get most of the way and miss the last third (the long-tail streams on hard tickets).

One more number, measured differently: averaged across tickets, a typical ticket has **90%** of its correct streams present in the examples — in plain counts, **~2.9 of its 3.2** correct Value Streams are found on average. (That 90% is a per-ticket *average*, not a count of tickets — a ticket that finds 3 of its 4 streams contributes 75% to it. Both numbers say the same thing: coverage is strong.)

2. Not every ticket is equally hard




How to read it. Each bar is *how many of the 373 tickets* have that many correct Value Streams:

- **Just 1 — 193 tickets (52%):** one right answer to find. Easy.
- **2-4 — 98, 5-9 — 54, 10+ — 28:** the more streams a ticket has, the harder it is to find them *all*. The 28 tickets with 10+ streams (a few have 19) are the truly hard cases.

So "90% coverage" is blended across these — the 52% single-stream tickets are easy and lift the average.

3. Hard tickets: we find most, but rarely all



How to read it. Two ticket groups (**easy** = 1 stream, 193 tickets; **hard** = 2+, 180 tickets), each measured two ways:


- **Blue** — "**Found MOST streams (avg)**": the average fraction of a ticket's streams that were found. Easy 92%, hard **88%** — nearly the same, so we find *most* of the answer even on hard tickets.
- **Green** — "**Found EVERY stream**": how often *all* streams were found. Easy 92%, hard only **63%**.

In plain counts:

- **Easy ticket:** has **1** correct Value Stream → found 92% of the time (**≈ 178 of 193** easy tickets get their one stream). With a single stream, "found most" and "found every" are the same — that's why the two bars match.
- **Hard ticket:** has **~5.6** correct Value Streams on average → the search finds **~4.9 of them** (88%), but lands the *complete* set only 63% of the time (**≈ 113 of 180** hard tickets).

So the hard-ticket gap (88% vs 63%) is the story: on a ~5.6-stream ticket we usually catch ~4.9 of them but **miss one or two from the long tail**, so we rarely get *all* of them.

4. Showing more tickets barely helps




How to read it. The x-axis is *how many past tickets we show* (6, 8, 10):

- **Green — Avg coverage:** climbs only 90% → 94% across the three ($\approx 2.9 \rightarrow 3.0$ of 3.2 streams found per ticket).
- **Blue — Found everything:** the share of tickets that get *all* their streams climbs only 78% → 85% ($\approx 291 \rightarrow 317$ of 373 tickets).

Each extra ticket adds a little coverage, but with **diminishing returns** — going from 6 to 10 examples finds only ~ 0.1 more stream on a typical ticket. **6 examples is enough** — more adds length without much benefit.

5. The right answer is usually at the very top



How to read it. For each ticket, the position of the *first* pulled ticket that carries a correct Value Stream:

- **Position 1 — 288 of 373 tickets (77%):** the very first result already carries a correct stream.
- **Positions 2, 3, 4–10 — ~76 tickets total:** a few need to look a little deeper.
- **None — 9 tickets (2%):** found nothing correct at all (the only true misses).

So the search not only *finds* the right streams, it usually *ranks them first* — for 3 in 4 tickets the top example is already on target.

All the numbers

Measure (per ticket)	Show 6	Show 8	Show 10
Streams found per ticket (of ~3.2 correct)	2.90	2.97	3.03
Avg coverage (fraction of streams found)	90%	92%	94%
Found at least one correct stream	96%	97%	98%
Found EVERY correct stream	78%	81%	85%

(A ticket has **3.2** correct Value Streams on average; the historic search surfaces **~2.9** of them. Note how the count barely moves — 2.90 → 3.03 — across 4 extra tickets: that's the diminishing returns from chart 4.)

Plain-language glossary

Term	What it means
Value Stream	The business category a ticket belongs to — what we predict.
Coverage	Of a ticket's correct Value Streams, how many showed up in the pulled examples.
Found everything	All of a ticket's correct Value Streams were found — none missed.
Easy / hard ticket	Easy = 1 correct Value Stream; hard = 2 or more.
@6 / @8 / @10	When showing the top 6 / 8 / 10 pulled past tickets.