


Can we feed the model the raw ticket text — and how big should we let it get?

Instead of an LLM-written summary, we want to give the model the raw ticket text (description + attachments). This checks whether that's feasible, and works out how many attachments to keep and what token budget to set. Measured on 374 tickets (a token $\approx \frac{3}{4}$ of a word; more tokens = slower and costlier).

Bottom line

- **Raw text is feasible** — most tickets are small, and even the largest fits the model.
 - **Most tickets are tiny:** the typical ticket's full text is **~3,854 tokens**. A few are huge (max ~89k), driven by big PowerPoint decks.
 - **18% of tickets have no attachments** — they run on the short description alone (median ~358 tokens), a low-context cohort worth flagging.
 - **Only ~half of attachments carry text** — of 1,579 attachments, 840 (53%) have extractable text; ~46% are images/unsupported. And **PowerPoint dominates by count + size but is the least text-dense** — **PDFs hold the real content**.
 - **The recommendation: a ~24k-token budget, packed greedily across attachments** — **no fixed count cap**. At ~24k, **95% of tickets fit untouched**; 16k (lean/fast, 86%) and 32k (generous, 98%) are the other operating points.
 - **Why budget, not a count:** the token budget is the real lever (a fixed attachment count barely changes coverage). So instead of "keep 4," we **pack greedily**: description first, then each attachment takes up to the remaining budget — small attachments don't waste their share, big ones aren't over-truncated, and the **5th/6th attachment is included whenever it still fits** (6-attachment tickets average ~24k, so they usually do).
-

1. How big is a ticket's text?




Each bar is the raw text size (description + all attachments) in tokens. **Typical (median) = 3,854** — half the tickets are smaller. The top 10% reach ~19k, the top 5% ~26k, and the single largest is ~89k (one giant deck). The description alone is tiny (median ~358 tokens) — **attachments are what make a ticket big, and only for a minority.**

The full distribution across all 374 tickets:

Most tickets' text is small; a long tail is large

142 of 374 tickets are under 2k tokens; 53 are over 16k.




142 of 374 tickets have under 2k tokens of text; the bulk are under 8k; only **53 tickets** are over 16k (the big-deck minority). The shape is a steep drop with a long right tail.

And how many tickets fit under each token budget (the raw text — description + all attachments):


What token budget fits most tickets?

Raw text (description + all attachments). 20k fits 91%; 25k fits 95%.



A 20k budget fits 91% of tickets; 25k fits 95%; 30k fits 97% (5k fits just 56%). This is the descriptive version of the budget decision (the cap-aware version we actually recommend is in section 7).

2. How many attachments does a ticket have?



Each bar is how many tickets have that many attachments:


- **0 — 67 tickets (18%):** no attachments; the model sees only the short description.
- **1–4 — 269 tickets (72%):** most tickets.
- **5+ — 38 tickets (10%):** the content-heavy minority.

3. How big is each attachment?

By text (tokens):

How big is each attachment (text)?

Most attachments are small (under 2k tokens); 75 are over 8k.




Most attachments are small — **290 carry under 1k tokens, 204 are 1–2k**. Only 75 attachments are over 8k tokens. So the typical attachment is light; a handful of big ones drive the tail.

By file size (on disk):


How big is each attachment (on disk)?

Many are small, but 226 attachments are over 2 MB (the decks).



Plenty of small files, but **226 attachments are over 2 MB** — the big PowerPoint decks. Median attachment is ~550 KB, but the largest is ~10 MB.

4. PowerPoint is most common, but PDF carries the most text




This is the standout finding. Comparing the three file types:

Type	How many	Avg text per file	File size	Text density
PowerPoint	410 (most common)	2,158 tokens	2.42 MB (biggest)	893 tokens/MB (lowest)
PDF	277	5,908 tokens (most)	1.16 MB	5,091 tokens/MB (highest)
Word	153	1,051 tokens	0.47 MB	2,243 tokens/MB

- **PowerPoint dominates by count and by disk size, but is the *least* text-dense** — decks are mostly images and layout, so a big 2.4 MB file yields only ~2.2k tokens.
- **PDFs carry the most actual text** (5.9k tokens/file, 6× the density of PowerPoint). The real content lives in the PDFs.

Implication: if attachments ever need trimming, dropping a big low-text PowerPoint costs little content; dropping a PDF costs a lot.


5. Only half of attachments actually carry text



Of 1,579 total attachments across the corpus:

- **840 (53%) carry extracted text** — the PDFs, PowerPoints, and Word docs.
 - **~726 (46%) are images or unsupported types** — no text at all (screenshots, diagrams, etc.).
 - **13 are empty** (supported but no extractable text).
 - **0 idea-card attachments were detected** — worth a flag: the "idea-card-first" selection rule never fires on this corpus (the filename pattern isn't matching), so condense always uses the ranked attachments packed into the budget.
-

6. More attachments → more tokens (and a jump at 5)



Average raw tokens for tickets with each attachment count (ticket count in parentheses):


- **Roughly linear up to 4** — each attachment adds **~3,000 tokens**: 1→3k, 2→6.7k, 3→9.5k, 4→11.7k.
- **A jump at 5** → **~26,520**: five-attachment tickets carry *bigger* attachments (~5.2k each), i.e. the big-deck tickets.

So **attachment count is a lever on the token size** — which is exactly why we choose the cap and the budget together (next).

7. The decision — budget is the real lever, not the cap

Token budget is the real lever — the attachment cap barely moves it


The lines sit close together: capping attachments hardly changes coverage (most tickets are small).



Each line is one attachment cap (keep 1 / 3 / 4 / all). The x-axis is the token budget; the y-axis is the % of tickets whose text fits under it (no truncation).

- **All lines rise steeply with budget** — 8k → 24k lifts coverage from ~68% to ~95%. **The budget moves the needle.**
- **The lines sit close together** (keep-4 and keep-all overlap) — **capping attachments hardly changes coverage**, because most tickets are small regardless; the cap only touches the ~10% with 5+.
- **Keep-1 is highest only because it throws away the most content** (next chart) — cheap coverage at a real cost.

8. Keeping 3–4 attachments retains ~all the content



How much of a ticket's attachment text survives each cap:

- **Keep 1 — 77%:** drops a quarter. Too aggressive.
- **Keep 3 — 97%, Keep 4 — 99%:** retains essentially everything — the knee.
- **Keep 5 / all — 100%:** the 5th-plus attachment adds under 3% — not worth the extra tokens.

So roughly **3–4 attachments' worth of budget keeps ~all the content** — which is why the greedy budget-packing in section 9 (which naturally fits ~that many before the budget bites) loses almost nothing, without needing a hard count cap.

9. The strategy: greedy budget-packing (no fixed count)

The budget sets coverage:

Budget	Tickets that fit	Truncated	Best for
12k	78%	22%	very tight latency
16k	86%	14%	lean / fast
24k	95%	5%	balanced (recommended)
32k	98%	2%	most generous

How the budget is applied — greedily, not as a fixed attachment count:

1. **Description** first (always, full) — it counts against the budget.

2. **Attachments** in ranked order (text-rich formats first): each takes up to the **remaining** budget; stop when the budget is exhausted.

This is better than the old "keep top-4, split the budget evenly": - a single big deck gets the room it needs instead of an arbitrary 1/N slice; - small attachments don't waste their slice; - the **5th/6th attachment is included when it still fits** (a 6-attachment ticket averages ~24k, so it usually does) — the old cap-4 would have thrown those away.

The **same budget applies to historical tickets** shown as precedent, so a giant past ticket can't blow up the prompt.

Implemented: `CondenseConfig.doc_char_budget = 96_000` (~24k tokens), greedy packing in `_consolidate`, `max_attachments = 8` (download/extract cap — the budget, not this count, caps content).

Why the count cap is 8 (and why it isn't the real limiter)

The 24k-token **budget** decides how much text survives — not the count. So why cap the count at all? Because the attachment ranker has to **download and extract** each candidate *before* the budget packer can decide what fits (and extraction is now a serialized native parse). Without a cap, a pathological ticket with 20–30 attachments would download and parse every one — wasting work on files that could never fit a 24k budget anyway. The cap bounds that wasted effort; it does **not** bound content.

Two EDA facts set the number at 8:

1. **The budget realistically holds ~4–6 attachments.** Each adds ~3k tokens up to 4 (1→3k, 2→6.7k, 3→9.5k, 4→11.7k); a 6-attachment ticket averages ~24k. So content-wise the budget is exhausted around the 5th–6th attachment.
2. **The cap must sit *above* that**, or it becomes the real limiter and re-introduces the old "keep top-4" truncation we deliberately removed. We want the **budget** to decide, with the count only catching outliers.

So **8 sits comfortably above the ~6 the budget can hold**: enough candidates for the greedy packer to fill 24k even when attachments are small, while still capping the rare 15-/20-/30-attachment tickets. Per the per-ticket distribution, the overwhelming majority of tickets have ≤ 8 attachments, so 8 almost never truncates a real set — it only bounds download/extract on outliers.

Honest caveat: 8 is a pragmatic guard, not a tuned optimum — anything in the ~6–10 range behaves almost identically because the budget does the real work. It only bites a ticket with many *small* attachments where the 9th/10th would have fit; the EDA says that's rare. If we ever want the budget to be the *sole* limiter, raise the cap to 10–12 or download lazily in ranked order until the budget fills.

The numbers

Token counts (per ticket)	Median	Top 10%	Top 5%	Max
Description only	358	880	1,241	2,634
Raw (description + all attachments)	3,854	19,433	25,921	88,614

Avg raw tokens by attachment count	0	1	2	3	4	5
avg tokens	552	3,042	6,708	9,489	11,652	26,520
tickets	67	93	66	56	54	22

Content kept by attachment cap	1	2	3	4	5
% of attachment text kept	77%	92%	97%	99%	100%

Coverage grid (% tickets that fit)	8k	12k	16k	24k	32k
keep 4 attachments	68%	78%	86%	95%	98%

Attachments	Value
Tickets with none	67 (18%)
Avg per ticket	2.25 (median 2, max 12)
Tokens per attachment	median 1,483 · max 68,768
File types	PowerPoint 49% · PDF 33% · Word 18%