

# Value Stream prediction — representation EDA

**Question:** for VS prediction, what text representation works best in each of the three places text is used — and can we drop summarization to save cost?

**Setup (held constant across every run):** - 100 tickets with  $\geq 3$  ground-truth VS ( `--min-gt 3 --sample 100` , fixed seed) — the hard, discriminating cases; single-VS tickets wash out differences. - `--count-mode gt` — request exactly the GT count, so micro **F1 = P = R** (one clean number per run) and no count-generosity bias. - evidence mode, the Recall selection prompt, K=6 historic neighbours, judge on.

## Terms & token counts (read this first)

Every IDMT ticket produces **two forms of text**. The whole study is about which form goes where.


| term                | what it actually is                                                                                                                                   | size                                                                 |
|---------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------|
| raw                 | the ticket's description + extracted attachment text, consolidated and capped during ingest                                                           | $\leq 24,000$ tokens<br>( <code>doc_char_budget = 96k</code> chars ) |
| summary             | the condense LLM's output from that raw — 6 fields: generatedSummary, businessProblem, businessCapability, keyTerms, stakeholders, systemsAndProducts | $\sim 460$ tokens (all 6 concatenated)                               |
| description         | only the ticket's Jira description field (no attachments)                                                                                             | a few hundred tokens                                                 |
| snippet             | the first $\sim 200$ characters of a stored doc                                                                                                       | $\sim 50$ tokens                                                     |
| <code>raw@Nk</code> | the <b>raw</b> truncated to N-thousand tokens (e.g. <code>raw@7k</code> = first $\sim 7,000$ tokens of the raw)                                       | N,000 tokens                                                         |

These forms get used in **three independent places** per prediction:

| place             | what it does                                                                | what goes here (winner)                                        |
|-------------------|-----------------------------------------------------------------------------|----------------------------------------------------------------|
| Retrieval query   | the text we <b>embed</b> to search the index for the 6 nearest past tickets | the <b>summary</b> (~460 tok)                                  |
| New-ticket prompt | the new ticket's text the LLM <b>reads to pick the VS</b>                   | the <b>full raw</b> (~24,000 tok)                              |
| Historic block    | each of the 6 retrieved neighbours, shown as precedent in the prompt        | each neighbour's <b>summary</b> (~460 tok × 6) + its VS labels |

So a run labelled "**raw + summary**" in the tables below means: **new-ticket prompt = raw**, **historic block = summary** (the retrieval query stays the summary unless the row says `raw@7k` retrieval). The ladder varies one place at a time.

## Results



Columns = the form used in each of the three places (see the terms table above).

| run         | new-ticket prompt | historic block | retrieval query | F1    | exact-set | avg latency |
|-------------|-------------------|----------------|-----------------|-------|-----------|-------------|
| all-summary | summary           | summary        | summary         | 0.715 | 23%       | 4.2s        |

| run                     | new-ticket prompt | historic block | retrieval query | F1           | exact-set  | avg latency |
|-------------------------|-------------------|----------------|-----------------|--------------|------------|-------------|
| <b>raw + summary</b>    | <b>raw</b>        | <b>summary</b> | <b>summary</b>  | <b>0.786</b> | <b>36%</b> | 5.8s        |
| raw + raw@1500          | raw               | raw@1500       | summary         | 0.781        | 26%        | 5.6s        |
| raw + raw@3000          | raw               | raw@3000       | summary         | 0.768        | 24%        | 6.3s        |
| raw + description       | raw               | description    | summary         | 0.780        | 31%        | 4.4s        |
| raw + raw@7k            | raw               | raw@7k         | summary         | 0.780        | 30%        | 9.4s        |
| <b>raw@7k retrieval</b> | raw@7k            | raw@7k         | <b>raw@7k</b>   | <b>0.742</b> | 23%        | 6.6s        |

The last row drops the summary entirely (raw-embedded index + raw everything). On a clean 100/100 with cheap raw@3k historic it lands at **0.742** — below the pack. (An earlier raw@7k-*historic* variant cost 13.9s avg / 132s max latency; dropping it to 3k fixed the latency but not the quality.)

**How to read it:** the big step is summary→raw on the **new-ticket prompt** (0.715 → 0.786). After that, swapping the **historic block** representation barely moves F1 (0.768–0.786) — and the last bar (raw@7k *retrieval*) drops *below* the pack. So the prompt is the lever, the historic block is a wash, and raw retrieval is a regression.


## Finding 1 — the lever is the NEW-TICKET prompt

Feeding the new ticket's **raw text** instead of its summary is **+0.071 F1 (0.715 → 0.786)** and nearly doubles exact-set (23% → 36%). Retrieval is already perfect at surfacing candidates (every GT reaches the LLM), so all of this gain is the LLM *choosing* better when it sees the full ticket.

## Finding 2 — the historic block representation is a wash (so use the cheapest)

Among the raw-prompt runs the historic block spans only 0.768–0.786: summary (0.786)  $\approx$  description (0.780)  $\approx$  raw@1500 (0.781)  $\approx$  raw@7k (0.780), with raw@3000 slightly *worse* (0.768). More raw precedent does **not** help — it dilutes. **Summary historic is marginally best and the cheapest.**

## Finding 3 — summary retrieval beats raw@7k retrieval (the cost question)




To test dropping summarization entirely, we re-embedded the index on **raw@7k** (no summary) and retrieved with raw. The decider is **historic-lane recall** — did the precedent search surface the GT (in evidence mode the 50-VS candidate pool makes review-pool recall structurally 1.0, so *that* isn't the signal — the historic lane is).

- summary retrieval: historic-lane R **0.902** → F1 **0.786**
- raw@7k retrieval: historic-lane R **0.843** → F1 **0.754**

**How to read it:** the ~460-token summary embeds the *whole* ticket cleanly; raw@7k keeps only ~half of a big ticket, so it retrieves **worse neighbours**. Since precedent drives recall (GT *backed* by historic is picked ~0.82, *not-backed* only ~0.38), worse precedent → lower recall → ~3 F1 points lost. **Dropping the summary costs quality.**

## Finding 4 — feed the new ticket's FULL raw; don't cap it




Holding everything else at the winner (summary retrieval + summary historic) and changing only the new ticket's raw cap: the **full ~24k raw** scores **0.780** vs **0.759** at a 7k cap — capping loses ~2 points and drops exact-set 31% → 26%. The extra context genuinely helps the LLM decide, and latency was identical (~4s) either way, so **there's no reason to cap the new-ticket prompt**.

## Where the time goes (latency split)

---


### Where the time goes (retrieval vs LLM)



**How to read it:** splitting prediction latency into retrieval vs the LLM call shows **retrieval is sub-second** (0.38–0.79s) — never the bottleneck. The **LLM call is the whole cost**, and it tracks the **historic block size**: summary historic 3.7s vs raw@3k historic 5.8s. So the cheap lever for latency is the historic representation (keep it summary), not the new-ticket prompt.

## Latency / cost

**Latency by representation (LLM prompt size drives it)**



**How to read it:** latency tracks the LLM **prompt size**. summary historic  $\approx 6 \times 460$  tokens; raw@7k historic  $\approx 6 \times 7k = 42k$  tokens, which blows the prompt to  $\sim 55k$  tokens and spikes prediction to 132s on big-neighbour tickets. This is a *second* cost axis, independent of the summary question: even if you keep summaries, **never ship raw@7k historic** — it's  $\sim 4\text{--}5\times$  the runtime token cost for zero quality.

## Verdict

**Locked config: summary retrieval + FULL  $\sim 24k$  raw new-ticket prompt + summary historic — F1  $\approx 0.78\text{--}0.79$ .**

- The summary you "can't eliminate" turns out to be the thing buying your best retrieval.
- Dropping summarization (raw@7k retrieval) is a **regression** ( $0.780 \rightarrow 0.742$ , historic-lane R  $0.902 \rightarrow 0.840$ ), not a marginal trade. The summary-free idea is dead.
- The new-ticket **raw** prompt is the real win ( $+0.07$ ) — and feed it **in full**; a 7k cap costs  $\sim 2$ pts.
- Use the **cheapest** historic block (summary): same quality as raw, and it's what keeps the LLM call (the only real latency) at  $\sim 3.7$ s. Retrieval is sub-second regardless.

## Final config (locked) — the per-ticket flow

1. New ticket raw, consolidated to  $\leq 24k$  tokens
2. Summarize it (condense LLM)  $\rightarrow$  summary fields
3. Embed the summary  $\rightarrow$  search the index  $\rightarrow$  retrieve the 6 nearest past tickets

4. Build the selection prompt:

- the new ticket's FULL raw (~24k) ("raw to decide")
- the 6 neighbours' summaries + their VS labels ("summaries as precedent")
- the 50-VS catalogue

5. LLM picks the VS (count = requested; evidence mode, Recall prompt)

The two uses of "summary" are different field sets, by design: - **Retrieval query** embeds **all six** summary fields (generatedSummary + businessProblem + businessCapability + keyTerms + stakeholders + systemsAndProducts). - **Historic block** shows each neighbour's **generatedSummary only** + its VS labels.

So every new ticket contributes **both** its summary (to *find* precedent) and its full raw (to *decide*) — which is exactly why summarization stays in the pipeline. **No further changes.**